

ECE 239 Week 3

01/26/2018

Tianwei Xing, Cheng Zheng

OH: Tue. 3pm-5pm, Tue. 10am-12pm

Eng. IV 67-112

Outline

- KNN
- Softmax
- SVM

KNN

Non-parametric method:


New instance is classified by a majority vote of its neighbors.

Pros & Cons

Distance:

Euclidean distance; (Text: Hamming)

Drawback: skewed dataset



KNN

Compute distance: L-n norm

Faster computation: Vectorization

- Reduce the nested for-loops (2 for-loops)
- Use broadcasting
- Expand $\|x_i - x_t\|^2 = (x_i - x_t)^T (x_i - x_t) = ?$

Predict labels: sort, find, vote

Find **optimal k and distance metric:** grid search?

Accuracy / Error-rate / N-fold cross validation

Softmax

A **generalization** of binary logistic regression classifier:

Produce a more intuitive output (normalized class prob)

$$\text{softmax}_i(\mathbf{x}) = \frac{e^{a_i(\mathbf{x})}}{\sum_{j=1}^c e^{a_j(\mathbf{x})}}$$

for $a_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_i$ and c being the number of classes.

If we let $\theta = \{\mathbf{w}_j, w_j\}_{j=1,\dots,c}$, then $\text{softmax}_i(\mathbf{x})$ can be interpreted as the probability that \mathbf{x} belongs to class i . That is,

$$\Pr(y^{(j)} = i | \mathbf{x}^{(j)}, \theta) = \text{softmax}_i(\mathbf{x}^{(j)})$$

Loss function: (cross-entropy loss)

$$\arg \min_{\theta} \sum_{i=1}^m \left(\log \sum_{j=1}^c e^{a_j(\mathbf{x})} - a_{y^{(i)}}(\mathbf{x}^{(i)}) \right)$$

Derived from MLE formulation: $\min -\log(\text{likelihood})$

Softmax

Numerical stability:

$$\log k = -\max_i a_i(\mathbf{x})$$

$$\begin{aligned}\text{softmax}_i(\mathbf{x}) &= \frac{e^{a_i(\mathbf{x})}}{\sum_{j=1}^c e^{a_j(\mathbf{x})}} \\ &= \frac{ke^{a_i(\mathbf{x})}}{k \sum_{j=1}^c e^{a_j(\mathbf{x})}} \\ &= \frac{e^{a_i(\mathbf{x}) + \log k}}{\sum_{j=1}^c e^{a_j(\mathbf{x}) + \log k}}\end{aligned}$$

Calculate the loss and grad: chain rule.

Vectorization: batch

Training: $W = W - \text{learning_rate} \times \text{grad}$

Prediction: top scoring class

Gradient descent


Hence, we arrive at gradient descent. To update \mathbf{x} so as to minimize $f(\mathbf{x})$, we repeatedly calculate:

$$\mathbf{x} := \mathbf{x} - \epsilon \nabla_x f(\mathbf{x})$$

ϵ is typically called the *learning rate*. It can change over iterations. Setting the value of ϵ appropriately is an important part of deep learning.

Optimizing softmax classifier: Using different learning rates to train classifiers and test their performance on validation data

Support vector machine



Hinge loss:

$$L_i = \sum_{j \neq y_i} [\max(0, x_i w_j - x_i w_{y_i} + \Delta)]$$

i iterates over all N samples,
 j iterates over all C classes,
 w_j is the weights for computing score of class j .
 y_i is the index of correct class of x_i
 $\Delta=1$ is the margin parameter.

Calculate the gradient

Covered in lecture, recall from lecture slides

Vectorization:

```
def fast_loss_and_grad(self, X, y):  
    """  
    A vectorized implementation of loss_and_grad. It shares the same  
    inputs and outputs as loss_and_grad.  
    """  
    loss = 0.0  
    grad = np.zeros(self.W.shape) # initialize the gradient as zero  
  
    # ===== #  
    # YOUR CODE HERE:  
    #   Calculate the SVM loss WITHOUT any for loops.  
    # ===== #  
  
    # ===== #  
    # END YOUR CODE HERE  
    # ===== #  
  
    # ===== #  
    # YOUR CODE HERE:  
    #   Calculate the SVM grad WITHOUT any for loops.  
    # ===== #  
  
    # ===== #  
    # END YOUR CODE HERE  
    # ===== #  
  
    return loss, grad
```

Hint: many of numpy function support
broadcasting

`numpy.maximum(x1, x2)`

Gradient Descent: the same as softmax