

In []:

```
import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized softmax loss. Store it as the variable loss.
        # (That is, calculate the sum of the losses of all the training
        # set margins, and then normalize the loss by the number of
        # training examples.)
        # ===== #

        num_examples = X.shape[0]
        num_classes = self.W.shape[0]
```

```

# print(num_examples, num_classes)

#outer summation
for i in range(0, num_examples):
    a_i = X[i].dot(self.W.T)

    #sum over all
    sigma_j = np.sum(np.exp(a_i))

    # -a_y(i)*X(i) + log(summation)
    probs = lambda idx: np.exp(a_i[idx])/sigma_j
    loss_sum = probs(y[i])#self.get_probs(y[i], a_i, sigma_j)

    #need to maintain negative sign
    loss -= np.log(loss_sum)

loss /= num_examples

# ===== #
# END YOUR CODE HERE
# ===== #

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
    the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #
    num_examples = X.shape[0]
    num_classes = self.W.shape[0]
# print(num_examples, num_classes)

#outer summation
"""

```

```

for i in range(0, num_examples):

    a_i = X[i].dot(self.W.T)

    #sum over all
    sigma_j = np.sum(np.exp(a_i))

    # -a_y(i)*X(i) + log(summation)
    probs = lambda idx: np.exp(a_i[idx])/sigma_j
    loss_sum = probs(y[i])

    #need to maintain negative sign
    loss -= np.log(loss_sum)

    #Gradient
    for c in range(0, num_classes):
        prob_class = probs(c)
        indicator = 0
        if(c == y[i]):
            indicator = 1

        grad_update = (prob_class - indicator)*X[i]

        grad[c, :] += grad_update

    """

for i in range(len(y)):
    score = self.W.dot(X[i,:].T)
    score -= np.max(score)
    true_score = score[y[i]]
    t_loss = np.exp(true_score) / np.sum(np.exp(score))
    loss -= np.log(t_loss)
    for j in range(self.W.shape[0]):
        indicator = 0
        if (j == y[i]):
            indicator = 1

        grad_update = (np.exp(score[j])/np.sum(np.exp(score)) - indicator)*X
[i]        grad[j, :] += grad_update

loss /= num_examples
grad /= num_examples

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """

```

```

sample a few random elements and only return numerical

in these dimensions.
"""

for i in np.arange(num_checks):
    ix = tuple([np.random.randint(m) for m in self.W.shape])

    oldval = self.W[ix]
    self.W[ix] = oldval + h # increment by h
    fxph = self.loss(X, y)
    self.W[ix] = oldval - h # decrement by h
    fxmh = self.loss(X,y) # evaluate f(x - h)
    self.W[ix] = oldval # reset

    grad_numerical = (fxph - fxmh) / (2 * h)
    grad_analytic = your_grad[ix]
    rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
    print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouptuts as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    num_classes = self.W.shape[0]
    num_features = self.W.shape[1]
    num_train = X.shape[0]

    scores = X.dot(self.W.T)
    scores_T = scores.T
    mask = np.zeros(shape = (num_classes, num_train))
    mask[y, range(num_train)] = 1

    gnd_truth = np.sum(np.multiply(mask, scores_T), axis=0)

    total_scores = np.sum(np.exp(scores_T).T, axis=1)
    log_inner = np.log(np.exp(gnd_truth) / total_scores)
    sigma = np.sum(log_inner, axis = 0)
    loss = -sigma
    loss /= num_train

    sigma_same_dims = np.sum(np.exp(scores), axis=1, keepdims=True)
    total_probs = np.exp(scores)/sigma_same_dims

```

```

total_probs[range(X.shape[0]), y] -= 1
grad = (X.T).dot(total_probs)
grad /= num_train
grad = grad.T

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (dim, batch_size)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with

```

```

# replacement.

# ===== #
rand_indices = np.random.choice(np.arange(num_train), batch_size)
X_batch = X[rand_indices]
y_batch = y[rand_indices]

# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate*grad

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    #   Predict the labels given the training data.
    # ===== #
    #y will be size N.
    # X=(N x D)  W=(C x D) where C = #of classes
    #Result will be (N x C)
    multi_class_preds = (X).dot(self.W.T)

    #find the highest ranking class value among the 10 classes -> columns so axis=1
    y_pred = np.argmax(multi_class_preds, axis=1)

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
  
return y_pred
```