

In []:

```
import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified
for ece239as at UCLA.
"""

class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized SVM loss, and store it as 'loss'.
        # (That is, calculate the sum of the losses of all the training
        # set margins, and then normalize the loss by the number of
        # training examples.)
        # ===== #
        for i in np.arange(num_train):
            predictions = X[i].dot(self.W.T)
            qnd truth = predictions[y[i]]
```

```

        counter = 0

    for j in range(0, num_classes):
        #j != y from summation
        if(j==y[i]):
            continue

        #margin calculation
        w = predictions[j] - gnd_truth + 1
        if(w > 0):
            counter+=1
            loss += w

    loss /= num_train # get mean

    return loss

def loss_and_grad(self, X, y):
    """
        Same as self.loss(X, y), except that it also returns the gradient.

        Output: grad -- a matrix of the same dimensions as W containing
                 the gradient of the loss with respect to W.
    """

    # compute the loss and the gradient

    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros_like(self.W)
    # print(self.W.shape[1])
    for i in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM loss and the gradient. Store the gradient in
        # the variable grad.
        # ===== #
        predictions = X[i].dot(self.W.T)
        gnd_truth = predictions[y[i]]
        counter = 0
        for j in range(0, num_classes):
            #j != y from summation
            if(j==y[i]):
                continue

            #margin calculation
            w = predictions[j] - gnd_truth + 1
            if(w > 0):
                counter+=1
                grad[j,:] += X[i]

```

```
loss += w
```

```
grad[y[i], :] += (-counter)*X[i]
```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #
```

```
loss /= num_train  
grad /= num_train
```

```
return loss, grad
```

```
def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
```

```
    """
```

```
    sample a few random elements and only return numerical  
    in these dimensions.
```

```
    """
```

```
    for i in np.arange(num_checks):
```

```
        ix = tuple([np.random.randint(m) for m in self.W.shape])
```

```
        oldval = self.W[ix]
```

```
        self.W[ix] = oldval + h # increment by h
```

```
        fxph = self.loss(X, y)
```

```
        self.W[ix] = oldval - h # decrement by h
```

```
        fxmh = self.loss(X,y) # evaluate f(x - h)
```

```
        self.W[ix] = oldval # reset
```

```
        grad_numerical = (fxph - fxmh) / (2 * h)
```

```
        grad_analytic = your_grad[ix]
```

```
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + a  
bs(grad_analytic))
```

```
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,  
grad_analytic, rel_error))
```

```
def fast_loss_and_grad(self, X, y):
```

```
    """
```

```
    A vectorized implementation of loss_and_grad. It shares the same  
    inputs and ouputs as loss_and_grad.
```

```
    """
```

```
    loss = 0.0
```

```
    grad = np.zeros(self.W.shape) # initialize the gradient as zero
```

```
    num_train = X.shape[0]
```

```
#     print(num_train)
```

```
#     print("X shape", X.shape)
```

```
# ===== #
```

```
# YOUR CODE HERE:
```

```
# Calculate the SVM loss WITHOUT any for loops.
```

```
# ===== #
```

```
predictions = X.dot(self.W.T)
```

```
gnd_truth = predictions[np.arange(num_train), y]
```

```

hinge = np.maximum(0, predictions - gnd_truth[:, np.newaxis] + 1)

hinge[np.arange(num_train), y] = 0
loss = np.sum(hinge)

loss /= num_train
# ===== #
# END YOUR CODE HERE
# ===== #

# ===== #
# YOUR CODE HERE:
# Calculate the SVM grad WITHOUT any for loops.
# ===== #

"""
-We have 3073 features, and 500 examples. There are 10 classes we can classifi
fy into.
-X is shape (500, 3073).
-The margins will be shape (500, 10), where each row corresponds to the marg
in for one
of the 500 examples. Index (i, j) corresponds to the margin of example i for
class j.
-We need to make a new matrix that has a 1 wherever the margin is > 0, and t
hen sum across the
columns.
"""
#Need to look at the margins and if > 0, we're going to need to add 1 to the
counter.
#Then, we multiply -counter by X

# print("margin shape", margins.shape)
counter_matrix = np.zeros(hinge.shape) #row = example, col = class

#place a 1 wherever hinge loss are > 0.
counter_matrix[hinge > 0] = 1

#Sum across the classes
counter = np.sum(counter_matrix, axis=1)

# print(counter)
# print(counter.shape)
# print(counter_matrix)

#Need to subtract the hinge loss values from each of the points in the count
er matrix
ex_idx = np.arange(num_train)
counter_matrix[ex_idx, y] = -counter #FIX THIS ERROR

#Take dot product of X and the errors to see how we need to update the weigh
ts
grad = (X.T.dot(counter_matrix)).T

```

```

grad /= num_train

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
          training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
          means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (dim, batch_size)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
        # replacement.
        # ===== #
        rand_indices = np.random.choice(np.arange(num_train), batch_size)
        # print(rand_indices)

```

```

X_batch = X[rand_indices]

y_batch = y[rand_indices]
# print(X.shape)
# print(X_batch.shape)
# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
# Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate*grad
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    N = the number of examples

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])

    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data with the parameter self.W.
    # ===== #
    #y will be size N.
    # X=(N x D) W=(C x D) where C = #of classes
    #Result will be (N x C)
    multi_class_preds = (X).dot(self.W.T)

    #find the highest ranking class value among the 10 classes -> columns so axis=1
    y_pred = np.argmax(multi_class_preds, axis=1)

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
  
return y_pred
```