

```

In [ ]: import numpy as np
import pdb

from .layers import *
from .layer_utils import *

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu). It has been modified in various areas for use
in the
ECE 239AS class at UCLA. This includes the descriptions of what code
to
implement as well as some slight potential changes in variable names to
be
consistent with class nomenclature. We thank Justin Johnson & Serena
Yeung for
permission to use this code. To see the original version, please visit
cs231n.stanford.edu.
"""

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden
    layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=0, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-2, dtype=np.float32, seed=None):
        """
        Initialize a new FullyConnectedNet.

```

```

    Inputs:
    - hidden_dims: A list of integers giving the size of each hidden layer.
    - input_dim: An integer giving the size of the input.
    - num_classes: An integer giving the number of classes to classify.
    - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
        the network should not use dropout at all.
    - use_batchnorm: Whether or not the network should use batch normalization.
    - reg: Scalar giving L2 regularization strength.
    - weight_scale: Scalar giving the standard deviation for random initialization of the weights.
    - dtype: A numpy datatype object; all computations will be performed using
        this datatype. float32 is faster but less accurate, so you should use
        float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout layers. This
        will make the dropout layers deterministic so we can gradient check the
        model.
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout > 0
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}

    # =====
#
# YOUR CODE HERE:
# Initialize all parameters of the network in the self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
#
# BATCHNORM: Initialize the gammas of each layer to 1 and the beta
# parameters to zero. The gamma and beta parameters for layer 1
# should be self.params['gamma1'] and self.params['beta1']. For layer

```

```

2, they
    # should be gamma2 and beta2, etc. Only use batchnorm if self.us
e_batchnorm
    # is true and DO NOT batch normalize the output scores.
    # =====e=====
= #
    mu = 0
    stddev = weight_scale
    """
    self.params['W1'] = std * np.random.randn(hidden_size, input_size)
    self.params['b1'] = np.zeros(hidden_size)
    self.params['W2'] = std * np.random.randn(output_size, hidden_size
)
    self.params['b2'] = np.zeros(output_size)

    np.random.normal(mu, stddev, <size>)
    """
    #aggregate all the dims into a single array that we can reference
    #input and output dim (num_classes) will only be used once
    aggregated_dims = [input_dim] + hidden_dims + [num_classes]
    for i in range(self.num_layers):
        #batchnorm on all layers except last one
        #init gammas to 1s and betas to 0
        if self.use_batchnorm and (i != (self.num_layers - 1)):
            self.params['gamma'+str(i+1)] = np.ones(aggregated_dims[i+1])
            self.params['beta'+str(i+1)] = np.zeros(aggregated_dims[i+1])

            self.params['b'+str(i+1)] = np.zeros(aggregated_dims[i+1])
            self.params['W'+str(i+1)] = np.random.normal(mu, stddev, size=(a
ggregated_dims[i], aggregated_dims[i+1]))
        # =====
#
        # END YOUR CODE HERE
        # =====
#

    # When using dropout we need to pass a dropout_param dictionary to
each
    # dropout layer so that the layer knows the dropout probability an
d the mode
    # (train / test). You can pass the same dropout_param to each drop
out layer.
    self.dropout_param = {}
    if self.use_dropout:
        self.dropout_param = {'mode': 'train', 'p': dropout}
        if seed is not None:
            self.dropout_param['seed'] = seed

    # With batch normalization we need to keep track of running means
and

```

```

    # variances, so we need to pass a special bn_param object to each
    batch
    # normalization layer. You should pass self.bn_params[0] to the fo
    rward pass
    # of the first batch normalization layer, self.bn_params[1] to the
    forward
    # pass of the second batch normalization layer, etc.
    self.bn_params = []
    if self.use_batchnorm:
        #for i in range(self.num_layers):
        self.bn_params = [{ 'mode': 'train'} for i in range(self.num_laye
rs - 1)]

    # Cast all parameters to the correct datatype
    for k, v in self.params.items():
        self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since
    they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None

    # =====
#
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    #
    # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm la
    yer
    # between the affine_forward and relu_forward layers. You may
    # also write an affine_batchnorm_relu() function in layer_utils.
    PY.
    #

```

```

# DROPOUT: If dropout is non-zero, insert a dropout layer after
# every ReLU layer.
# =====
#

nn_layer = {}
nn_cache = {}
batchnorm_cache = {}
dropout_cache = {}

#initialize the first layer with the inputs
nn_layer[0] = X
#pass through each layer
for i in range(1, self.num_layers):
#    print("iteration", i)
    gamma_idx = 'gamma'+str(i)
    beta_idx = 'beta'+str(i)
    w_idx = 'W'+str(i)
    b_idx = 'b'+str(i)
    #affine relu forward takes (x, w, b)
    if self.use_batchnorm:
        #args: x, gamma, beta, bn_param
        nn_layer[i], batchnorm_cache[i] = batchnorm_forward(nn_layer[
i-1], self.params['gamma'+str(i)], self.params['beta'+str(i)], self.bn
_params[i-1])
        #    print(nn_layer[i-1].shape, self.params[w_idx].shape)
        nn_layer[i], nn_cache[i] = affine_batchnorm_relu_forward(nn_la
yer[i-1], self.params[w_idx],
                                                                    self
.params[b_idx], self.params[gamma_idx],
                                                                    self
.params[beta_idx], self.bn_params[i-1])
    else:
        nn_layer[i], nn_cache[i] = affine_relu_forward(nn_layer[i-1],
self.params[w_idx], self.params[b_idx])

    if(self.use_dropout):
        nn_layer[i], dropout_cache[i] = dropout_forward(nn_layer[i], s
elf.dropout_param)
    #all layers will have the affine_relu except for the last layer, w
hich is a passthrough
    #affine_forward takes (x, w, b) and outputs out, cache
    w_idx = 'W'+str(self.num_layers)
    b_idx = 'b'+str(self.num_layers)
    scores, cached_scores = affine_forward(nn_layer[self.num_layers -1
], self.params[w_idx], self.params[b_idx])

    nn_cache[self.num_layers] = cached_scores
# =====

```

```

#
# END YOUR CODE HERE
# =====
#

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# =====
#
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
#
# BATCHNORM: Incorporate the backward pass of the batchnorm.
#
# DROPOUT: Incorporate the backward pass of dropout.
# =====
#

#get loss w/ softmax loss
loss, grad_loss = softmax_loss(scores, y)

#add L2 regularization to loss 1/2*np.sum(w**2)
for i in range(1, self.num_layers + 1):
    cur_weight_matrix = self.params['W'+str(i)]
    loss += 0.5 * self.reg * np.sum(cur_weight_matrix**2)

    """
    Backpropping into the (n-1)th layer will be different because we don't have
    the relu. Use affine_backward and then for each previous layer apply affine_relu_backward
    affine_backward takes dout, cache and returns dx, dw, db
    affine_relu_backward takes dout, cache and returns dx, dw, db
    """

    dx={}
    w_idx_nth = 'W'+str(self.num_layers)
    b_idx_nth = 'b'+str(self.num_layers)
    dx[self.num_layers], grads[w_idx_nth], grads[b_idx_nth] = affine_backward(grad_loss, cached_scores)

# print(dx[3])
#regularize
grads[w_idx_nth] += self.reg * self.params[w_idx_nth]

```

```

#we apply affine_relu_backward now
for i in range(self.num_layers - 1, 0, -1):
#     print(i, self.num_layers)

    #dx, dw, db
    w_idx = 'W' + str(i)#+1)
    b_idx = 'b' + str(i)#+1)

    gamma_idx = 'gamma' + str(i)#+1)
    beta_idx = 'beta' + str(i)#+1)

    if self.use_dropout:
        dx[i+1] = dropout_backward(dx[i+1], dropout_cache[i])
    if self.use_batchnorm:
        dx[i], grads[w_idx], grads[b_idx], grads[gamma_idx], grads[bet
a_idx] = affine_batchnorm_relu_backward(dx[i+1], nn_cache[i])

    else:
        #dout input to affine_relu_backward is the
        dx[i], grads[w_idx], grads[b_idx] = affine_relu_backward( dx[i
+1], nn_cache[i])

        #regularize
#     print(grads[w_idx].shape, self.params[w_idx].shape )
    grads[w_idx] += self.reg * self.params[w_idx]

    # =====
#
# END YOUR CODE HERE
# =====
#
return loss, grads

```