

```

In [ ]: import numpy as np
        from nn1.layers import *
        import pdb

        """
        This code was originally written for CS 231n at Stanford University
        (cs231n.stanford.edu). It has been modified in various areas for use
        in the
        ECE 239AS class at UCLA. This includes the descriptions of what code
        to
        implement as well as some slight potential changes in variable names t
        o be
        consistent with class nomenclature. We thank Justin Johnson & Serena
        Yeung for
        permission to use this code. To see the original version, please visi
        t
        cs231n.stanford.edu.
        """

        def conv_forward_naive(x, w, b, conv_param):
            """
            A naive implementation of the forward pass for a convolutional layer
            .

            The input consists of N data points, each with C channels, height H
            and width
            W. We convolve each input with F different filters, where each filte
            r spans
            all C channels and has height HH and width HH.

            Input:
            - x: Input data of shape (N, C, H, W)
            - w: Filter weights of shape (F, C, HH, WW)
            - b: Biases, of shape (F,)
            - conv_param: A dictionary with the following keys:
              - 'stride': The number of pixels between adjacent receptive fields
            in the
              horizontal and vertical directions.
              - 'pad': The number of pixels that will be used to zero-pad the in
            put.

            Returns a tuple of:
            - out: Output data, of shape (N, F, H', W') where H' and W' are give
            n by
               $H' = 1 + (H + 2 * pad - HH) / stride$ 
               $W' = 1 + (W + 2 * pad - WW) / stride$ 
            - cache: (x, w, b, conv_param)
            """

```

```

out = None
pad = conv_param['pad']
stride = conv_param['stride']

# ===== #
# YOUR CODE HERE:
#   Implement the forward pass of a convolutional neural network.
#   Store the output as 'out'.
#   Hint: to pad the array, you can use the function np.pad.
# ===== #
N, C, H, W = x.shape
F, C2, HH, WW = w.shape

H_prime = 1 + (H + 2*pad - HH)/stride
W_prime = 1 + (W + 2*pad - WW)/stride

#pad the input
x_pad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad, pad)), mode='constant', constant_values=0)
_, _, H_pad, W_pad = x_pad.shape

out = np.zeros((N, F, int(H_prime), int(W_prime)))

#iterate through all data points
for datapoint in range(N):
    x_pad_cur = x_pad[datapoint]

    h_loc, w_loc = -1, -1
    #go by height
    for hi in range(0, H_pad - HH + 1, stride):
        h_loc += 1
        #go by width
        for wi in range(0, W_pad - WW + 1, stride):
            w_loc += 1
            #first dim = : to get all channels
            x_all_channels = x_pad_cur[:, hi:hi+HH, wi:wi+WW]

            for filt in range(F):
                out[datapoint, filt, h_loc, w_loc] = np.sum(x_all_channels * w[filt]) + b[filt]

            #reset height counter
            w_loc = -1

# ===== #
# END YOUR CODE HERE
# ===== #

cache = (x, w, b, conv_param)

```

```

    return out, cache

def conv_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a convolutional layer.

    Inputs:
    - dout: Upstream derivatives.
    - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive

    Returns a tuple of:
    - dx: Gradient with respect to x
    - dw: Gradient with respect to w
    - db: Gradient with respect to b
    """
    dx, dw, db = None, None, None

    N, F, out_height, out_width = dout.shape
    x, w, b, conv_param = cache

    stride, pad = [conv_param['stride'], conv_param['pad']]
    xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
    num_filts, _, f_height, f_width = w.shape

    # ===== #
    # YOUR CODE HERE:
    # Implement the backward pass of a convolutional neural network.
    # Calculate the gradients: dx, dw, and db.
    # ===== #

    #b is biases of shape (F,)
    #dout is N, F, out_height, out_width
    db = np.zeros((b.shape))
    for i in range(F):
        db[i] = np.sum(dout[:, i, :, :])

    #w: Filter weights of shape (F, C, HH, WW)
    F, C, HH, WW = w.shape
    dw = np.zeros((w.shape))
    for i in range(F):
        for j in range(C):
            for k in range(HH):
                for l in range(WW):
                    #Input data of shape (N, C, H, W)
                    derivative = dout[:, i, :, :] * xpad[:, j, k:k + out_height
                    * stride:stride, l:l + out_width * stride:stride]

```

```

        dw[i,j,k,l] = np.sum(derivative)

#x: (N, C, H, W)
_, _, H, W = x.shape
#create dummy gradient -> will have same dimensions as x
dx = np.zeros(x.shape)
#pad the gradient dx
dxdpad = np.pad(dx, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')

H_prime = 1 + (H + 2*pad - HH)/stride
W_prime = 1 + (W + 2*pad - WW)/stride

for i in range(N): #for each data point
    for j in range(F):
        for k in range(0, int(H_prime)): #, stride):
            k_prime = k*stride
            for l in range(0, int(W_prime)): #, stride):
                l_prime = l*stride
                #multiply the weights of this filter by derivative dout
                derivative = w[j] * dout[i,j,k,l]
            #
            # print(dxdpad.shape)
            # print(derivative.shape)
            dxdpad[i, :, k_prime:k_prime + HH, l_prime:l_prime+WW] += derivative

#extract derivative
#dimensions need to be pulled out from H, W (, , H, W)
dx = dxdpad[:, :, pad:pad+H, pad:pad+W]
# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def max_pool_forward_naive(x, pool_param):
    """
    A naive implementation of the forward pass for a max pooling layer.

    Inputs:
    - x: Input data, of shape (N, C, H, W)
    - pool_param: dictionary with the following keys:
      - 'pool_height': The height of each pooling region
      - 'pool_width': The width of each pooling region
      - 'stride': The distance between adjacent pooling regions

    Returns a tuple of:
    - out: Output data
    - cache: (x, pool_param)

```

```

"""
out = None

# ===== #
# YOUR CODE HERE:
#   Implement the max pooling forward pass.
# ===== #
pool_height = pool_param['pool_height']
pool_width = pool_param['pool_width']
stride = pool_param['stride']

N, C, H, W = x.shape
W_out = (W - pool_width)/stride + 1
H_out = (H - pool_height)/stride + 1

out = np.zeros((N, C, int(H_out), int(W_out)))

for datapoint in range(N):
    #reduce by one dimension (datapoint #)
    x_cur = x[datapoint]

    h_loc, w_loc = -1, -1
    for hi in range(0, H-pool_height + 1, stride):
        h_loc += 1
        for wi in range(0, W-pool_width + 1, stride):
            w_loc += 1

            #this is the receptive field
            x_receptive_field = x_cur[:, hi:hi+pool_height, wi:wi+pool_width]

            #iterate through all channels
            for c in range(C):
                out[datapoint, c, h_loc, w_loc] = np.max(x_receptive_field[c, :, :])

            w_loc = -1

    h_loc = -1

# ===== #
# END YOUR CODE HERE
# ===== #
cache = (x, pool_param)
return out, cache

def max_pool_backward_naive(dout, cache):
    """

```

```

A naive implementation of the backward pass for a max pooling layer.

Inputs:
- dout: Upstream derivatives
- cache: A tuple of (x, pool_param) as in the forward pass.

Returns:
- dx: Gradient with respect to x
"""
dx = None
x, pool_param = cache
pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']

# ===== #
# YOUR CODE HERE:
# Implement the max pooling backward pass.
# ===== #

dx = np.zeros(x.shape)
N, C, H, W = x.shape
H_prime = 1 + (H - pool_height)/stride
W_prime = 1 + (W - pool_width)/stride
for i in range(N):
    for j in range(C):
        for k in range(int(H_prime)):
            for l in range(int(W_prime)):
                k_prime = k * stride
                l_prime = l * stride

                #we want to only reward for the one we picked
                cur_window = x[i, j, k_prime:k_prime + pool_height, l_prime:l_prime + pool_width]
                max_cur_window = np.max(cur_window)
                masked_window = (cur_window == max_cur_window)
                derivative = dout[i, j, k, l] * masked_window
                dx[i, j, k_prime:k_prime + pool_height, l_prime:l_prime + pool_width] += derivative

# ===== #
# END YOUR CODE HERE
# ===== #

return dx

def spatial_batchnorm_forward(x, gamma, beta, bn_param):
    """
    Computes the forward pass for spatial batch normalization.

    Inputs:
    - x: Input data of shape (N, C, H, W)

```

```

- gamma: Scale parameter, of shape (C,)
- beta: Shift parameter, of shape (C,)
- bn_param: Dictionary with the following keys:
  - mode: 'train' or 'test'; required
  - eps: Constant for numeric stability
  - momentum: Constant for running mean / variance. momentum=0 means
that
    old information is discarded completely at every time step, whil
e
    momentum=1 means that new information is never incorporated. The
    default of momentum=0.9 should work well in most situations.
  - running_mean: Array of shape (D,) giving running mean of feature
s
  - running_var Array of shape (D,) giving running variance of featu
res

Returns a tuple of:
- out: Output data, of shape (N, C, H, W)
- cache: Values needed for the backward pass
"""
out, cache = None, None

# ===== #
# YOUR CODE HERE:
#   Implement the spatial batchnorm forward pass.
#
#   You may find it useful to use the batchnorm forward pass you
#   implemented in HW #4.
# ===== #
#mode = bn_param['mode']
# eps = bn_param.get['eps']
# momentum = bn_param.get['momentum']
N, C, H, W = x.shape

#reshape the (N, C, H, W) array as an (N*H*W, C) array and perform b
atch normalization on this array.
transpose = np.transpose(x, axes=(0,2,3,1))
reshaped = transpose.reshape(N*H*W, C)
bn_out, cache = batchnorm_forward(reshaped, gamma, beta, bn_param)

#reshape again and swap
out = bn_out.reshape(N, H, W, C).transpose(0, 3, 1, 2)

# ===== #
# END YOUR CODE HERE
# ===== #

return out, cache

```

```

def spatial_batchnorm_backward(dout, cache):
    """
    Computes the backward pass for spatial batch normalization.

    Inputs:
    - dout: Upstream derivatives, of shape (N, C, H, W)
    - cache: Values from the forward pass

    Returns a tuple of:
    - dx: Gradient with respect to inputs, of shape (N, C, H, W)
    - dgamma: Gradient with respect to scale parameter, of shape (C,)
    - dbeta: Gradient with respect to shift parameter, of shape (C,)
    """
    dx, dgamma, dbeta = None, None, None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the spatial batchnorm backward pass.
    #
    #   You may find it useful to use the batchnorm forward pass you
    #   implemented in HW #4.
    # ===== #

    N, C, H, W = dout.shape
    transpose = np.transpose(dout, axes=(0,2,3,1))
    reshaped = transpose.reshape(N*H*W, C)
    dx_bn, dgamma_bn, dbeta_bn = batchnorm_backward(reshaped, cache)

    dx = dx_bn.reshape(N, H, W, C).transpose(0,3,1,2)
    dgamma = dgamma_bn
    dbeta = dbeta_bn

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dx, dgamma, dbeta

```